



livingsocial

don't fear the
threads

simplify your life with JRuby

david.copeland@livingsocial.com

@davetron5000

www.naildrivin5.com

www.awesomelineapps.com



```
{  
:me => {  
:tech_lead =>  
“LivingSocial”  
}  
}
```

First Assignment

Write a

Write a
Command
Line
App

Build Awesome Command-Line Applications in Ruby

Control Your Computer,
Simplify Your Life



David Bryant Copeland

Edited by John Osborn

The Facets of Ruby Series



Write a
Command
Line
App

Write a Command Line App

To charge
credit cards
faster

**NO NEW
HARDWARE**

THREADS!





What problem
are we solving?

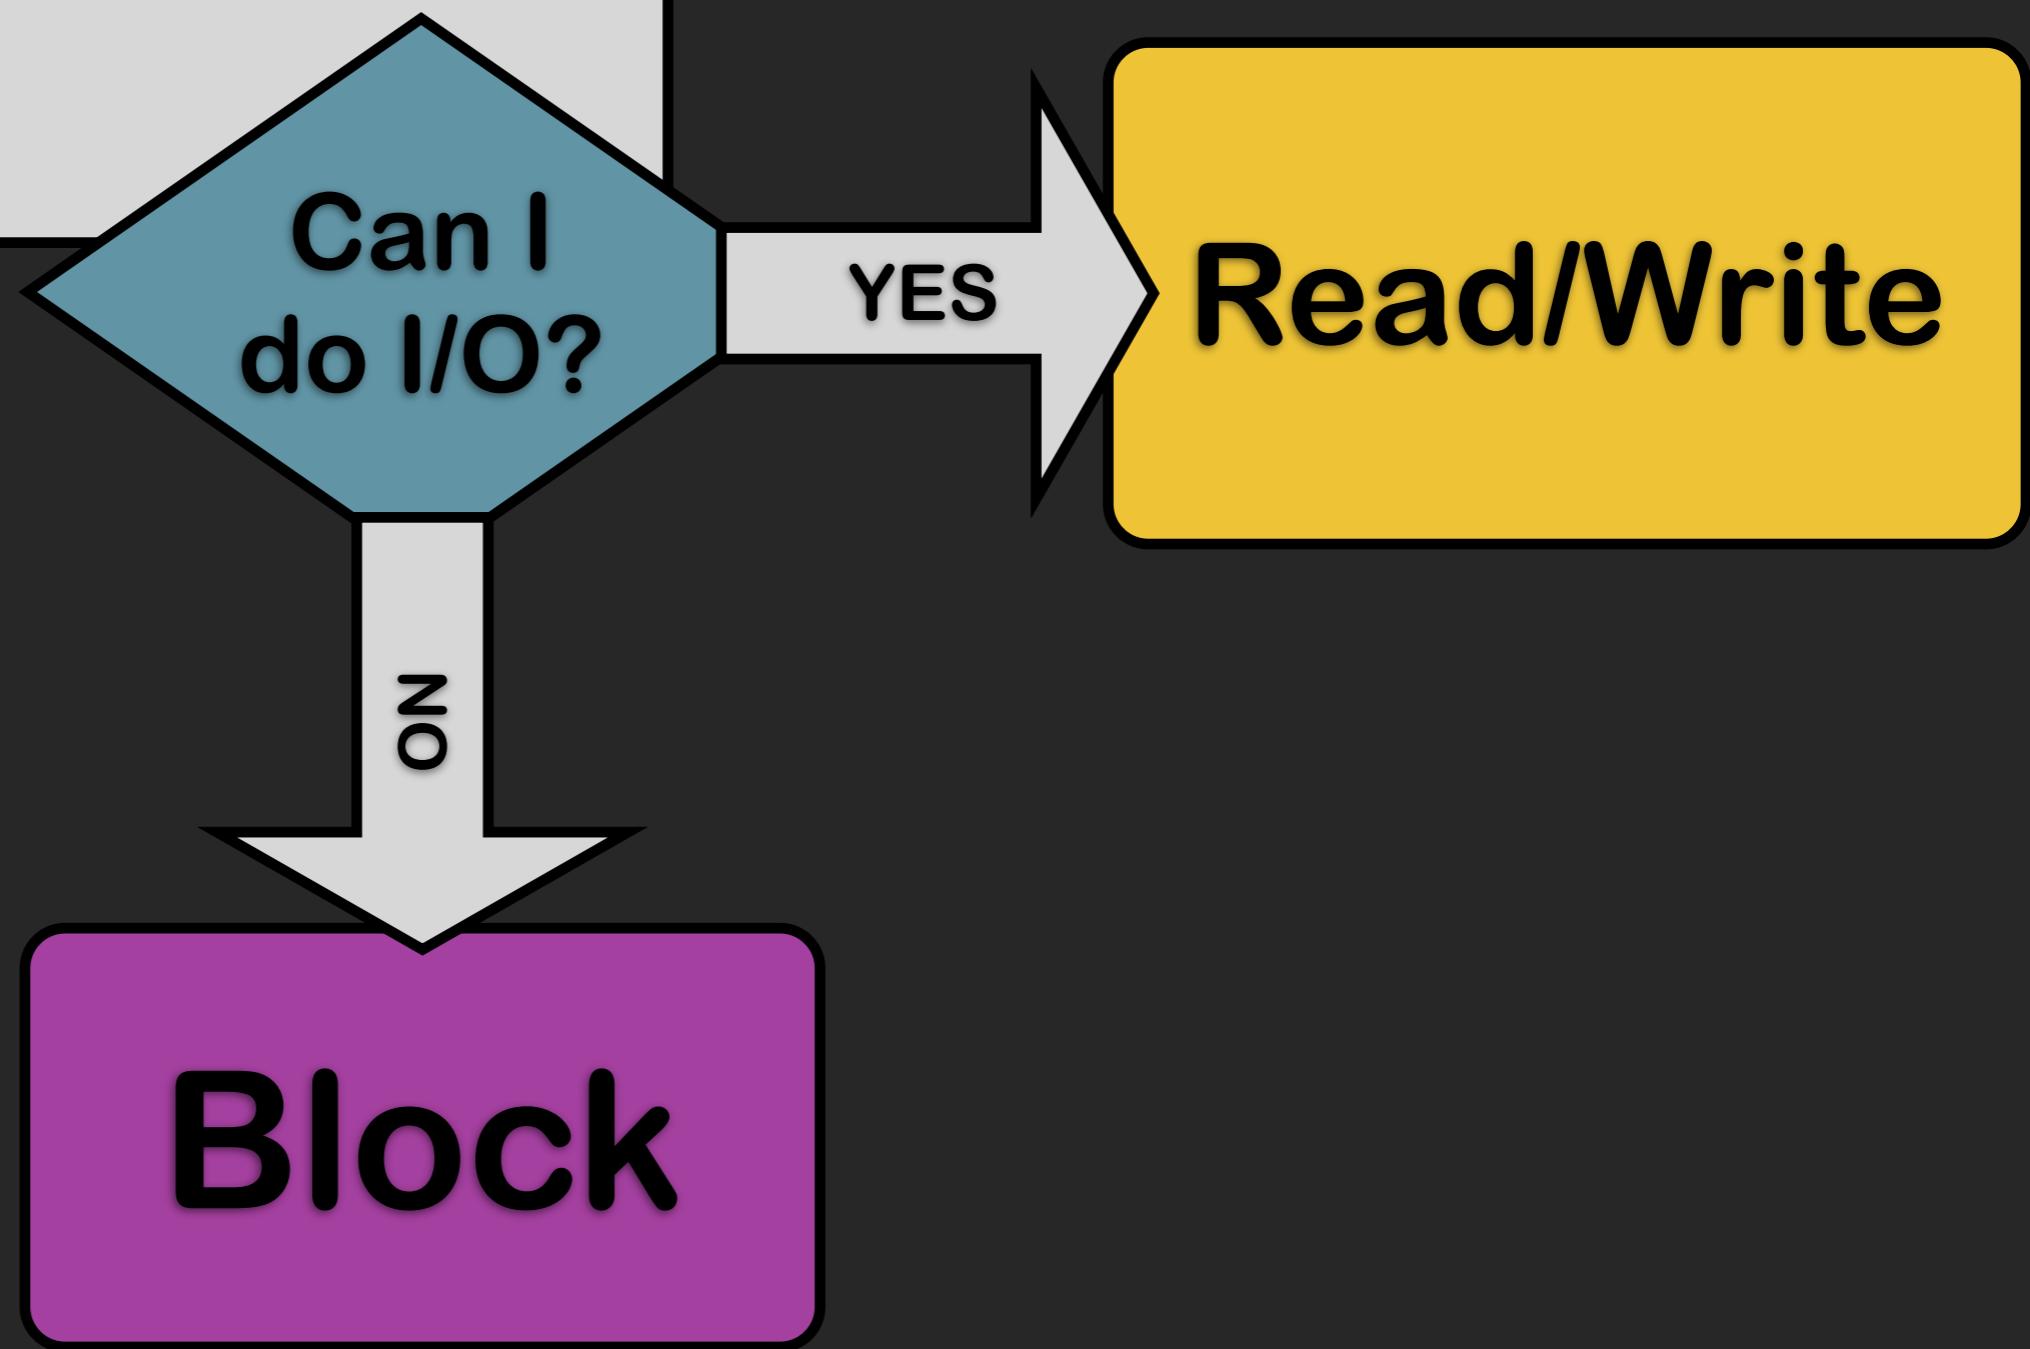
Do some I/O

**Compute
Stuff**



Do some I/O

Do some I/O



Do some I/O

Block

Block

Block

Let someone
else work

Block

**Maximize
Resources**

will our
code
be

Will our
code
be easy to write?

Will our
code
be easy to understand?
easy to write?

Will our
code
be easy to test?
easy to understand?
easy to write?

Are we
maximizing
our
resources?

Run lots of
processes


```
fork {  
    # your code  
}
```

easy to
understand

doesn't
maximize resources

Parent

**Parent's
Memory**

Parent

**Parent's
Memory**

Parent

fork

**Parent's
Memory**

Child

(copy of)

Event-based I/O


```
require 'em-http'
include EM::HttpRequest
```

```
EM.run {
```

```
# some of your code that does IO
some_callback { |results_of_io|
```

```
# the rest of your code that
# uses the results
```

```
}
```

```
}
```

```
require 'em-http'  
include EM::HttpRequest
```

```
EM.run {  
  # some of your code that does IO  
  some_callback { |results_of_io|  
    # the rest of your code that  
    # uses the results  
  }  
}
```

Need to do
I/O again?


```
require 'em-http'  
include EM::HttpRequest
```

```
EM.run {
```

```
  # some of your code that does SOME of your I/O  
  some_callback { |results_of_io|
```

```
    # use the results
```

```
    # now some more code that does some MORE I/O
```

```
    some_other_callback { |results_of_more_io|
```

```
      # use the results of THIS I/O
```

```
}
```

```
}
```

```
}
```



```
require 'em-http'
include EM::HttpRequest

EM.run {

    # some of your code that does SOME of your I/O
    some_callback { |results_of_io|
        # use the results
        # now some more code that does some MORE I/O
        some_other_callback { |results_of_more_io|
            # use the results of THIS I/O

            # SERIOUSLY, more I/O?!??!@ How much do you need?
            yet_more_callbacks { |moar_resultz|
                # How many levels deep are we now!
            }
        }
    }
}
```

```
require 'em-http'  
include EM::HttpRequest
```

```
EM.run {
```

```
# some of your code that does SOME of your I/O  
some_callback { |results_of_io|  
    # use the results  
    # now some more code that does some MORE I/O  
    some_other_callback { |results_of_more_io|  
        # use the results of THIS I/O
```

```
# SERIOUSLY, more I/O?!??!@ How much do you need?  
yet_more_callbacks { |moar_resultz|  
    # How many levels deep are we now?
```

```
}
```

NOT easy to
understand

kinda
maximizes resources

Entire call
chain must
be evented

Not true
parallelism

Not true
parallelism
unless you run lots
of processes

What's
parallelism?

Concurrency

Parallelism

Parallelism

Code running at the same time
as other code

Parallelism

Impossible with only one CPU

Code running at the same time
as other code

Event-based I/O

NOT easy to
understand

kinda
maximizes resources

Threads


```
Thread.new {  
  # your code  
}
```

```
fork {  
    # your code  
}
```

```
Thread.new {  
  # your code  
}
```

Need to do
I/O?

Need to do
I/O?

Not a problem

Need to do
I/O THREE
TIMES?

Need to do
I/O THREE
TIMES?

Not a problem

What if I need
to calculate π
to the
2,345,123rd
decimal
place?

**NOT A
PROBLEM**

Threads achieve
true
parallelism

easy to
understand

maximizes
resources

Threads
don't work in
Ruby

Threads
don't work in
Ruby...right?

MRI (aka C Ruby)

MRI (aka C Ruby)

Thread is an OS thread
(except in 1.8)

MRI (aka C Ruby)

Thread is an OS thread
(except in 1.8)

GIL :(

MRI (aka C Ruby)

Thread is an OS thread
(except in 1.8)

GIL :(

I/O will cause a context switch

Rubinius

Rubinius

Thread is an OS thread

Rubinius

Thread is an OS thread

No GIL!

Rubinius

Thread is an OS thread

No GIL!

True parallelism

JRuby

JRuby

Thread is a JVM Thread
(which is an OS thread)

JRuby

Thread is a JVM Thread
(which is an OS thread)

Context switch for variety of
reasons

JRuby

Thread is a JVM Thread
(which is an OS thread)

Context switch for variety of
reasons

True parallelism

JRuby

BATTLE TESTED

So, you've
decided to
use
Threads....

You need to
know four
things

Start & Manage

Start & Manage

(Dealing with) Shared State

Start & Manage

(Dealing with) Shared State

Using Third Party Libraries

Start & Manage (Dealing with) Shared State Using Third Party Libraries Context Switching

Start & Manage Threads

Chaos

```
Thread.new {  
    # your code  
}  
  
Thread.new {  
    # moar code  
}  
  
# etc
```

Hand-Roll

```
threads = []
threads << Thread.new {
    # your code
}
threads << Thread.new {
    # moar code
}
# etc
threads.each(&:join)
# All threads have completed
exit 0
```

`java.util.concurrent`

```
import java.util.concurrent # JRuby only

service =
  Executors.new_fixed_thread_pool(100)

service.execute {
  # your code
}
service.execute {
  # some other code
}
```

Method Summary

<code>static Callable<Object> callable(PrivilegedAction<?> action)</code>	Returns a <code>Callable</code> object that, when called, performs the specified action.
<code>static Callable<Object> callable(PrivilegedExceptionAction<?> action)</code>	Returns a <code>Callable</code> object that, when called, performs the specified action and handles any exception.
<code>static Callable<Object> callable(Runnable task)</code>	Returns a <code>Callable</code> object that, when called, performs the specified task.
<code>static <T> Callable<T> callable(Runnable task, T result)</code>	Returns a <code>Callable</code> object that, when called, performs the specified task and returns the specified result.
<code>static ThreadFactory defaultThreadFactory()</code>	Returns a default thread factory used to create threads.
<code>static ExecutorService newCachedThreadPool()</code>	Creates a thread pool that creates new threads as needed.
<code>static ExecutorService newCachedThreadPool(ThreadFactory threadFactory)</code>	Creates a thread pool that creates new threads as needed using the specified <code>ThreadFactory</code> to create new threads when needed.
<code>static ExecutorService newFixedThreadPool(int nThreads)</code>	Creates a thread pool that reuses a fixed number of threads.
<code>static ExecutorService newFixedThreadPool(int nThreads, ThreadFactory threadFactory)</code>	Creates a thread pool that reuses a fixed number of threads using the specified <code>ThreadFactory</code> to create new threads when needed.

Method Summary

static <code>Callable<Object></code>	<code>callable(PrivilegedAction<?> action)</code> Returns a <code>Callable</code> object that, when called, performs the specified action.
static <code>Callable<Object></code>	<code>callable(PrivilegedExceptionAction<?> action)</code> Returns a <code>Callable</code> object that, when called, performs the specified exception action.
static <code>Callable<Object></code>	<code>callable(Runnable task)</code> Returns a <code>Callable</code> object that, when called, performs the specified task.
static <code><T> Callable<T></code>	<code>callable(Runnable task, T result)</code> Returns a <code>Callable</code> object that, when called, performs the specified task and returns the specified result.
static <code>ThreadFactory</code>	<code>defaultThreadFactory()</code> Returns a default thread factory used to create threads.
static <code>ExecutorService</code>	<code>newCachedThreadPool()</code> Creates a thread pool that creates new threads as needed.
static <code>ExecutorService</code>	<code>newCachedThreadPool(ThreadFactory threadFactory)</code> Creates a thread pool that creates new threads as needed using the specified ThreadFactory to create new threads when needed.
static <code>ExecutorService</code>	<code>newFixedThreadPool(int nThreads)</code> Creates a thread pool that reuses a fixed number of threads.
static <code>ExecutorService</code>	<code>newFixedThreadPool(int nThreads, ThreadFactory threadFactory)</code> Creates a thread pool that reuses a fixed number of threads using the specified ThreadFactory to create new threads when needed.

Method Summary

static Callable<Object>

callable(PrivilegedAction<?> action)

Returns a Callable object that, when cal-

static Callable<Object>

callable(PrivilegedExceptionAction<?>

Returns a Callable object that, when cal-

static Callable<Object>

callable(Runnable task)

Returns a Callable object that, when cal-

<T> static Callable<T>

callable(Runnable task, T result)

Returns a Callable object that, when cal-

static ThreadFactory

defaultThreadFactory(ThreadFactory threadFactory)

Returns a default ThreadFactory used to cre-

static ExecutorService

newCachedThreadPool()

Creates a thread pool that creates new thr-

static ExecutorService

newCachedThreadPool(ThreadFactory threadFactory)

Creates a thread pool that creates new thr-

ThreadFactory to create new threads when need-

static ExecutorService

newFixedThreadPool(int nThreads)

Creates a thread pool that reuses a fixed num-

static ExecutorService

newFixedThreadPool(int nThreads, ThreadFactory threadFactory)

Creates a thread pool that reuses a fixed num-
when needed.

ExecutorService

ExecutorService

ExecutorService

execute()

ExecutorService

execute()

shutdown()

ExecutorService

`execute()`

`shutdown()`

`await_termination()`

ExecutorService

`execute()`

`shutdown()`

`await_termination()`

`shutdown_now()`

```
service = Executors.new_fixed_thread_pool(10)
tcp_server = TCPServer.new("127.0.0.1", 8080)

loop do {
  s = tcp_server.accept
  service.execute {
    s.puts calculate_pi()
    s.close
  }
}
```

```
service = Executors.new_fixed_thread_pool(10)
tcp_server = TCPServer.new("127.0.0.1", 8080)

Signal trap('SIGINT') {
    service.shutdown
}

loop do {
    s = tcp_server.accept
    service.execute {
        s.puts calculate_pi()
        s.close
    }
    break if service.is_shutdown
}
```

```
service = Executors.new_fixed_thread_pool(10)
tcp_server = TCPServer.new("127.0.0.1", 8080)

Signal trap('SIGINT') {
    service.shutdown
}

loop do {
    s = tcp_server.accept
    service.execute {
        s.puts calculate_pi()
        s.close
    }
    break if service.is_shutdown
}
service.await_termination(10, TimeUnit.SECONDS)
service.shutdown_now
```

So much more

So much more

ScheduledExecutorService

So much more

ThreadFactory

ScheduledExecutorService

So much more

Read the
javadocs

ThreadFactory

ScheduledExecutorService

Shared State

```
results = []
Thread.new {
    results << some_result()
}

Thread.new {
    results << other_result()
}
```

```
mutex = Mutex.new

results = []
Thread.new {
  mutex.synchronize {
    results << some_result()
  }
}
Thread.new {
  mutex.synchronize {
    results << other_result()
  }
}
```

Avoid explicit
locking/sync

```
results = ConcurrentLinkedQueue.new
```

```
Thread.new {  
    results.offer(some_result())  
}
```

```
Thread.new {  
    results.offer(some_result())  
}
```

```
import java.util.concurrent.atomic  
  
flag = AtomicBoolean.new(true)  
Thread.new {  
    if flag.get {  
        flag.set(false)  
    }  
}  
Thread.new {  
    flag.set(true)  
}
```

So much more

So much more

ConcurrentHashMap

So much more

AtomicReference

ConcurrentHashMap

So much more

Read the
javadocs

AtomicReference

ConcurrentHashMap

Third Party Libraries

Are they
thread safe?

Probably

Probably
...but to be sure

variables

Global variables

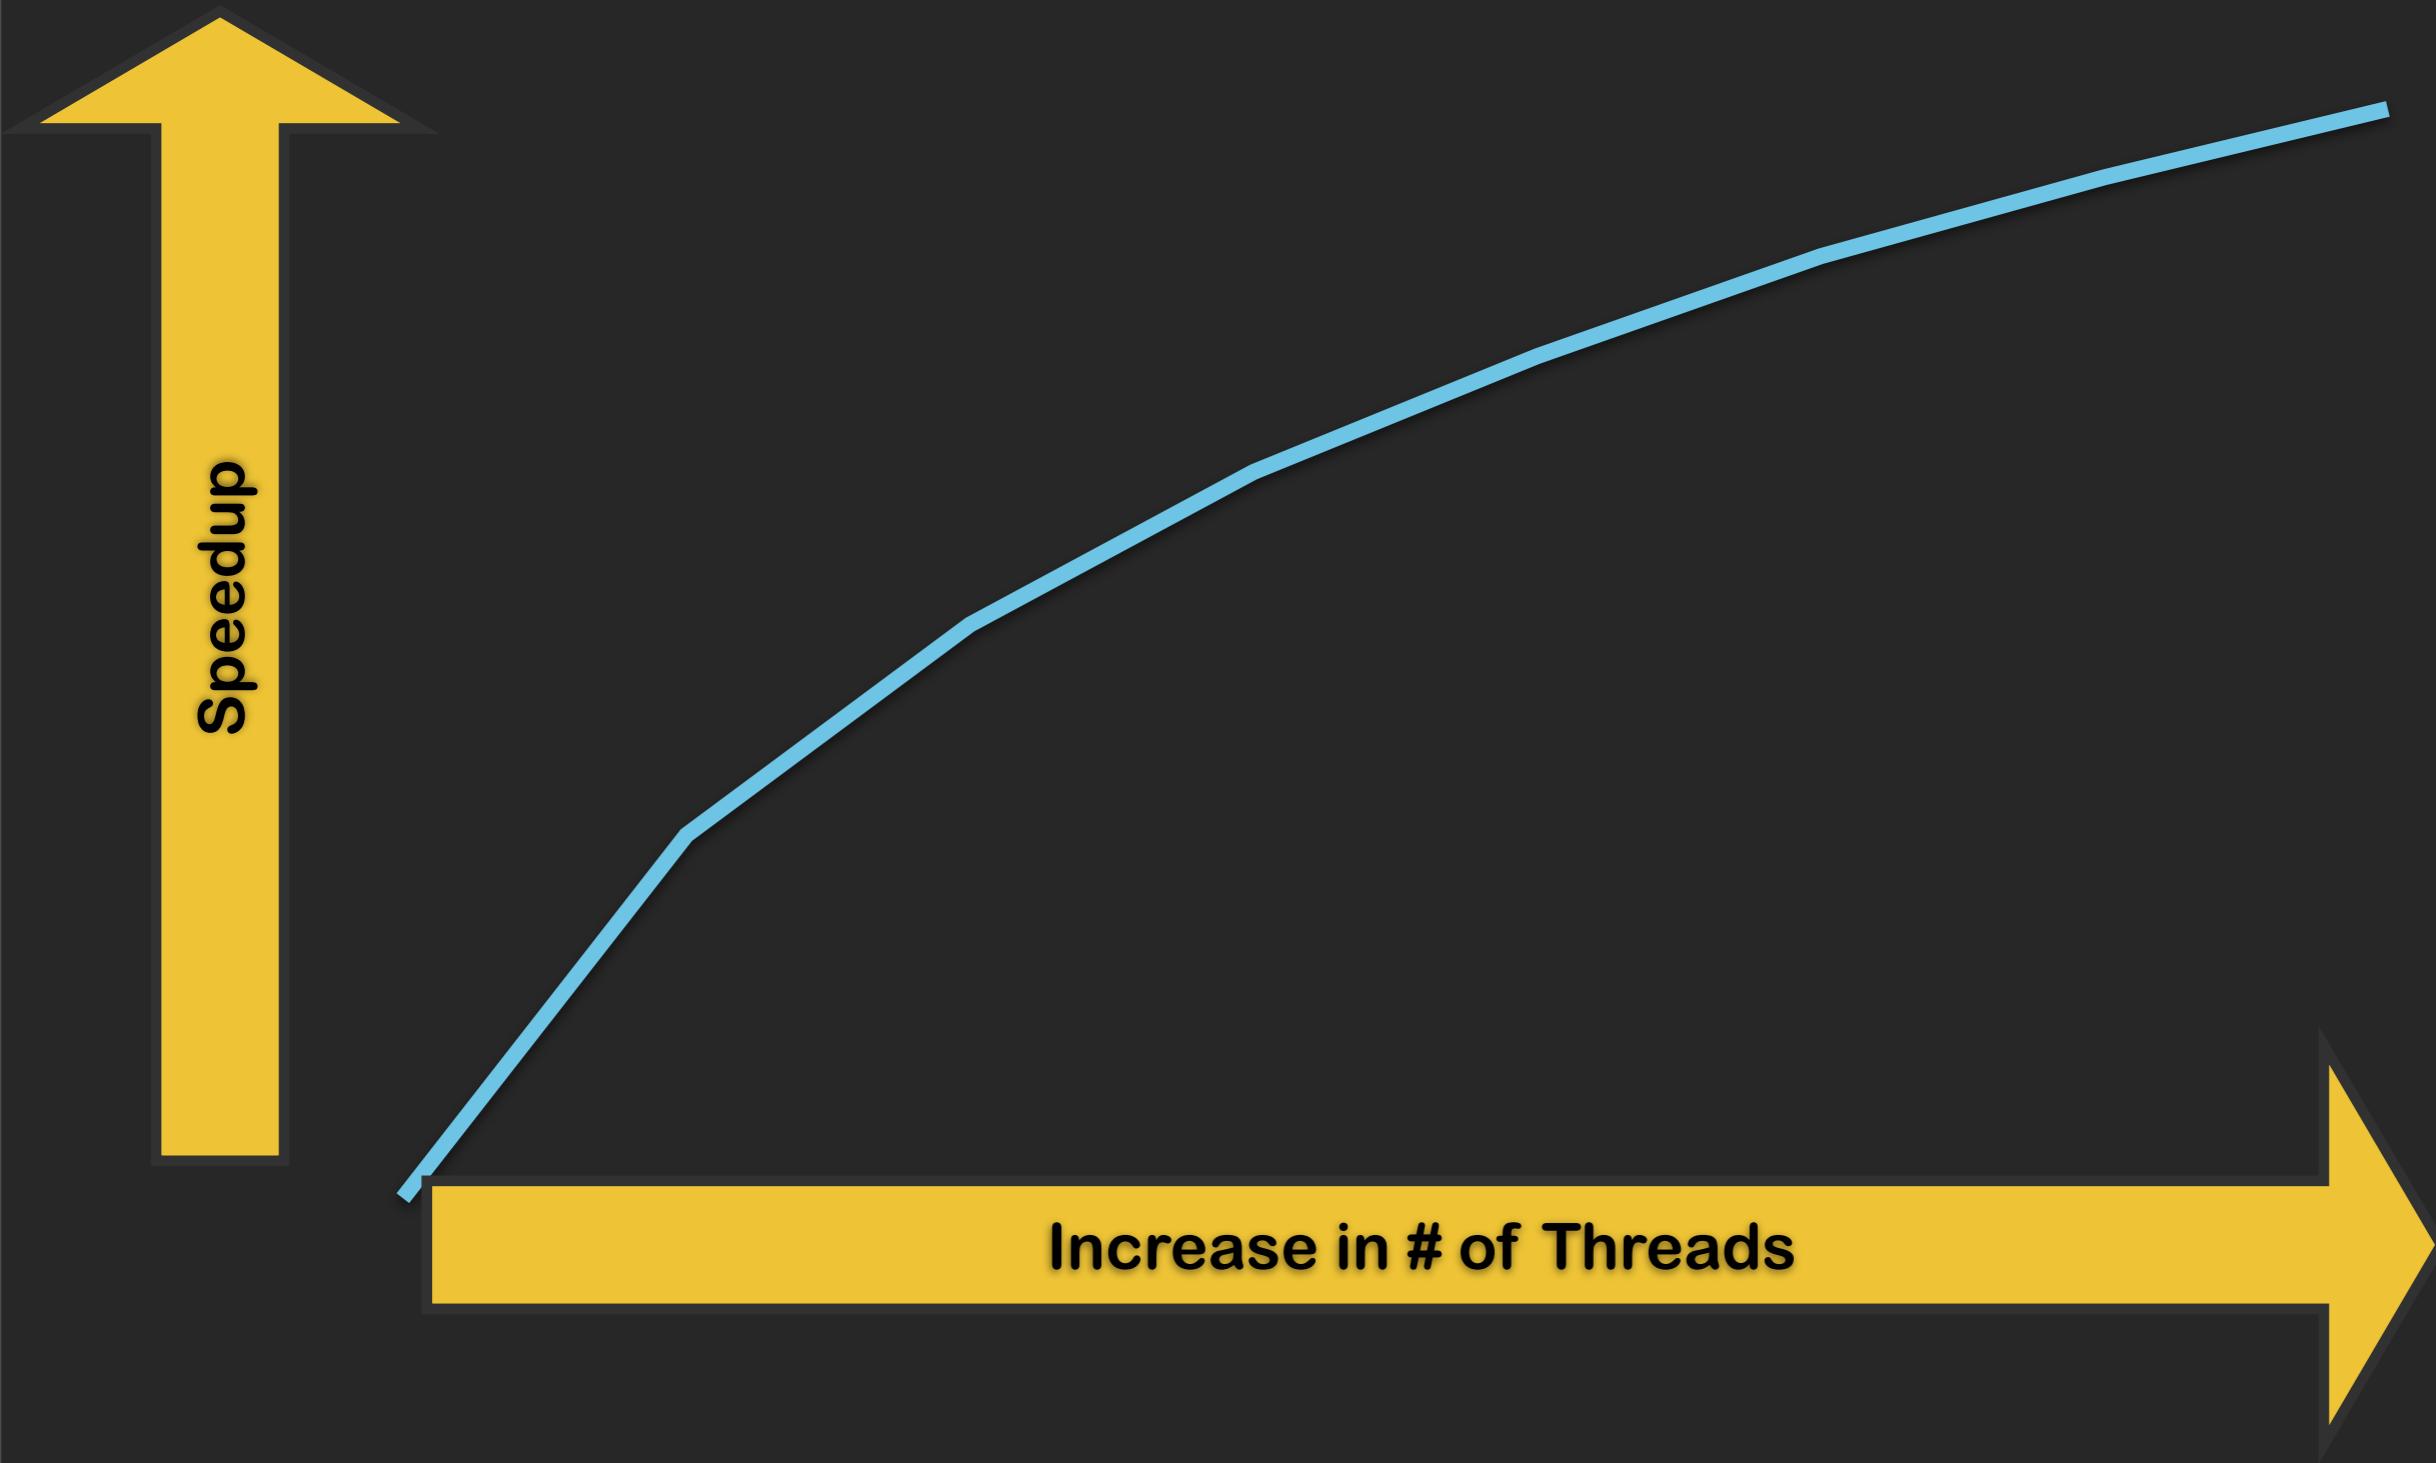
Class variables
Global

Most are fine

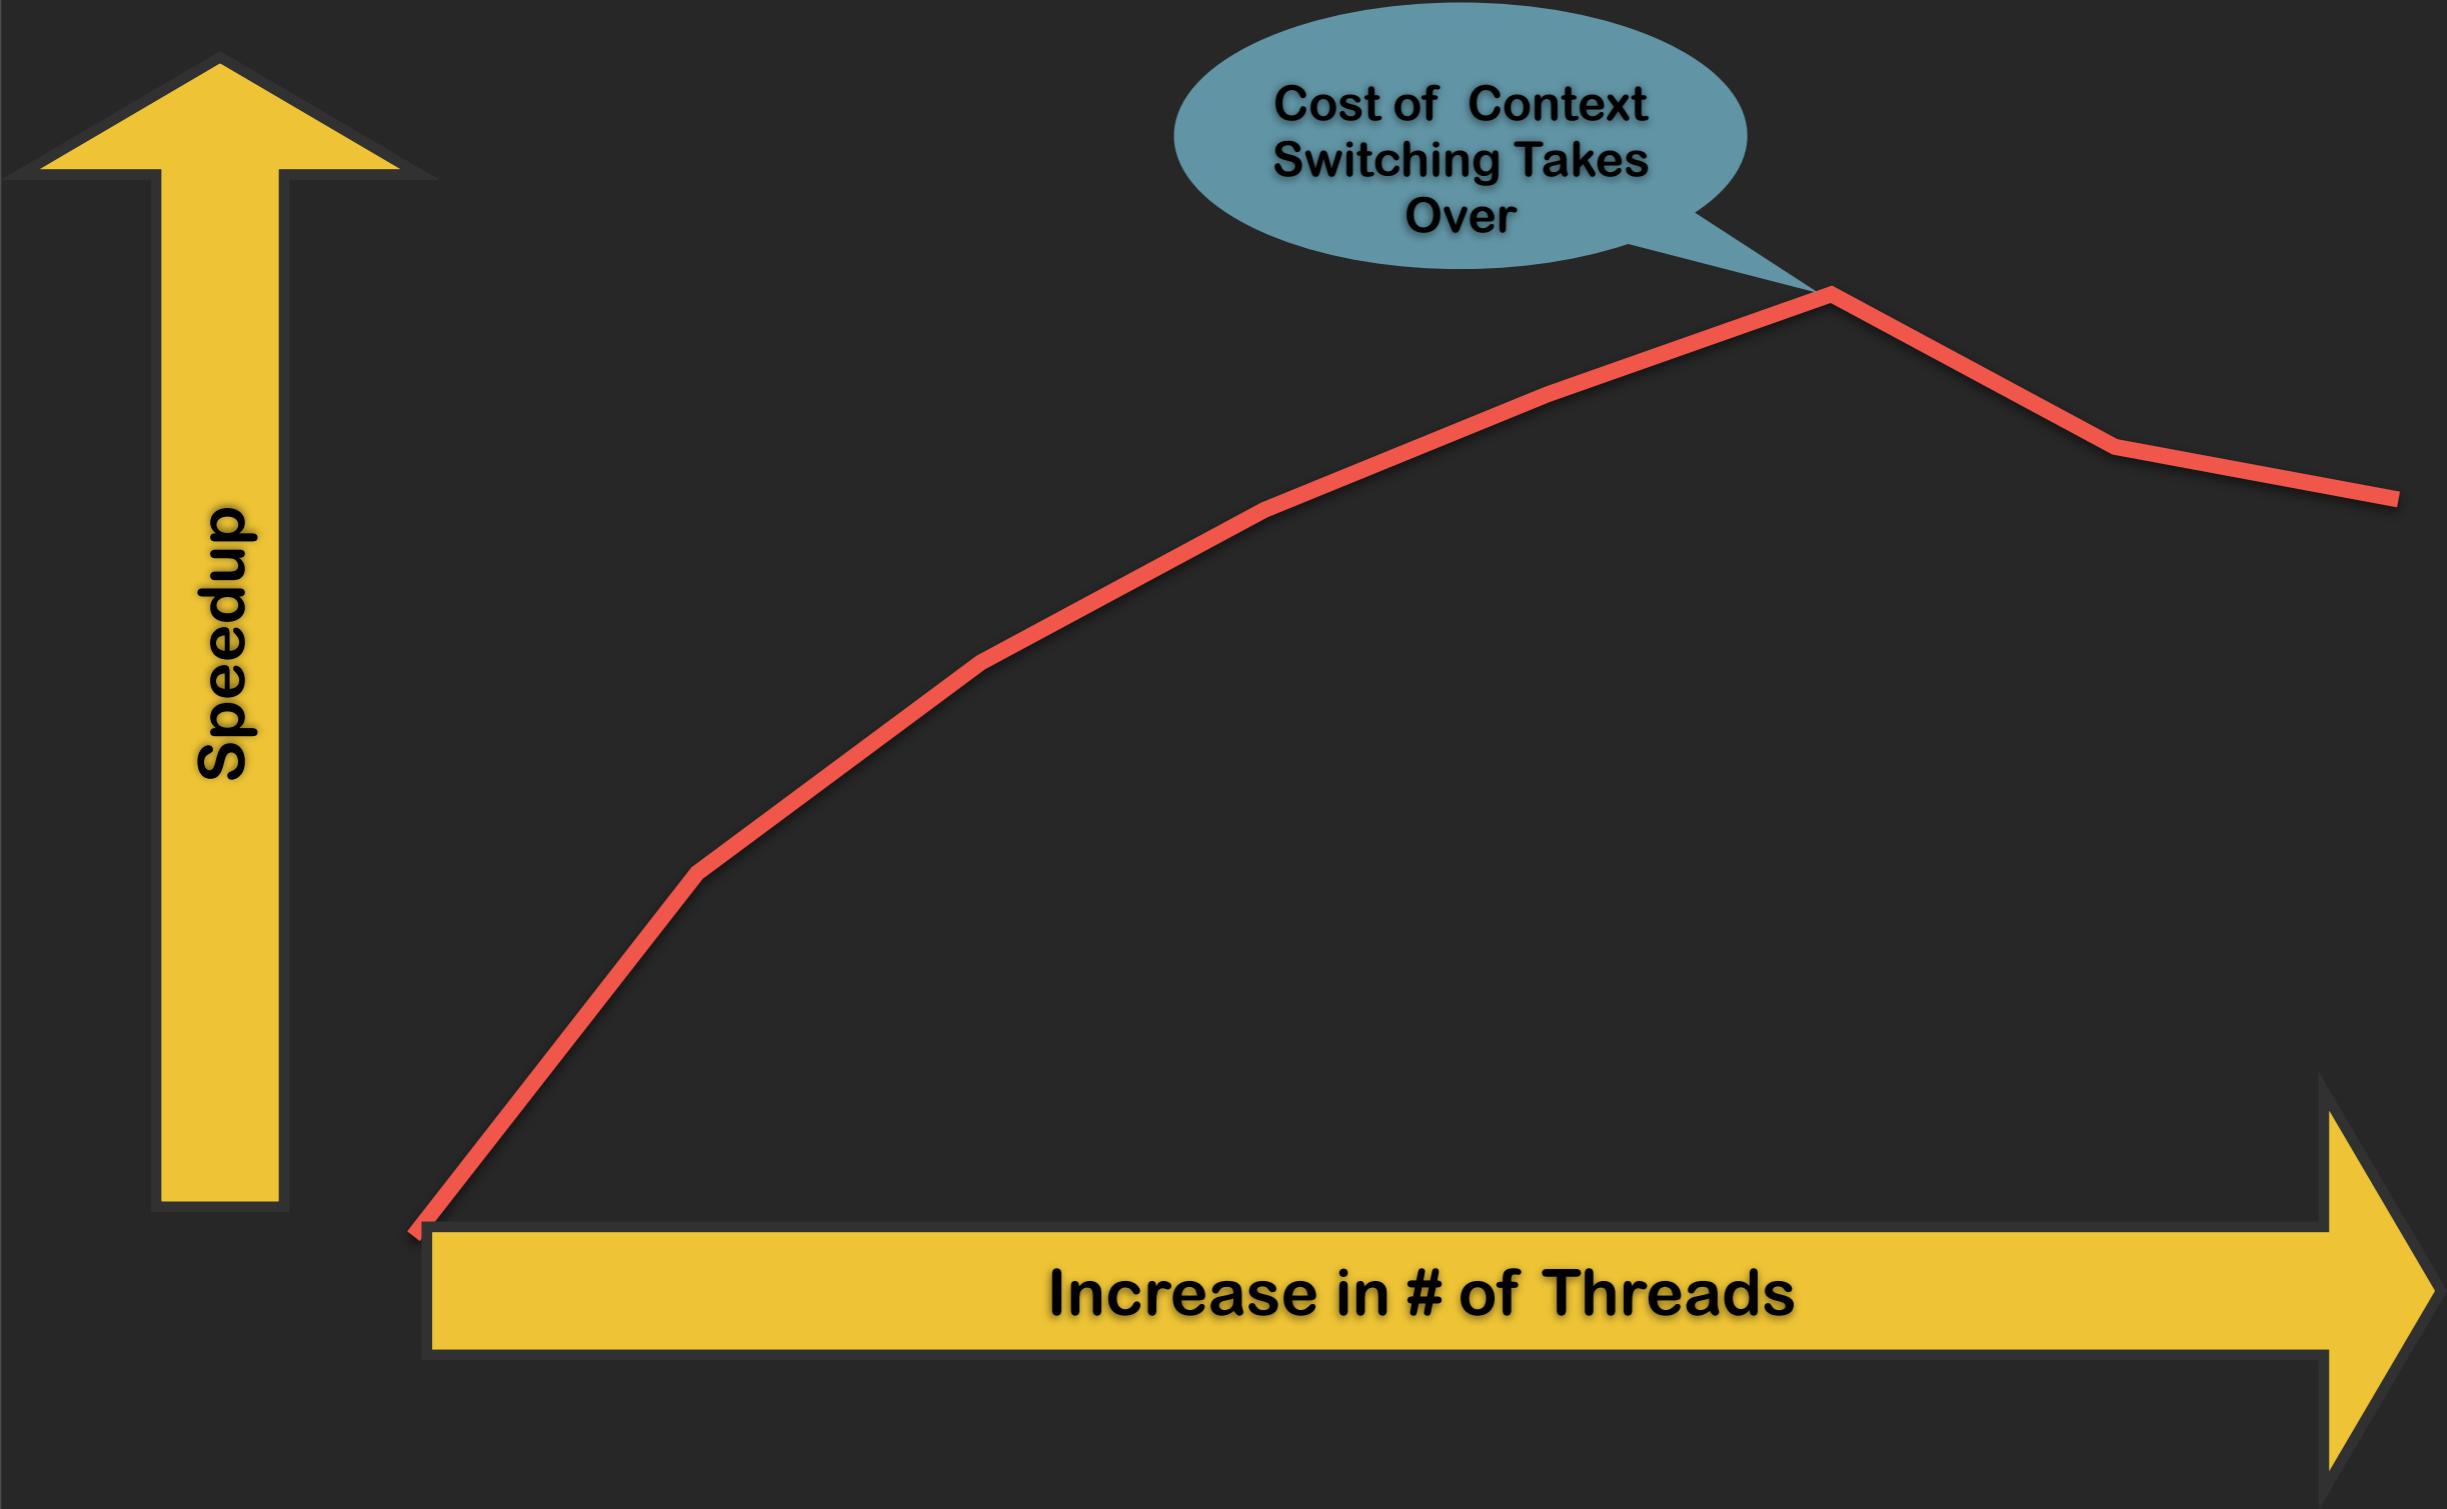
Context Switching

Amdahl's Law

— Performance Gains



— Performance Gains with Context Switching



Thread Pools

Worker

Worker

Worker

Worker

Worker

Worker

Worker

Worker

Thread Pool

Thread

Thread

Thread

Worker

Worker

Worker

Worker

Worker

Thread Pool

**Thread
Worker**

**Thread
Worker**

**Thread
Worker**

Worker

Worker

Worker

Worker

Thread Pool

Thread
Worker

Thread
Worker

Thread
Worker

```
Thread.new {  
  # your code  
}
```

```
service.execute {  
    # your code  
}
```

Not usually a
concern

Always use
Threads, right?

Thread UNsafe Libraries

Evented

Simple, I/O
bound task

Evented

Cannot use
JRuby

Evented

Otherwise,
Threads will
simplify your code
and maximize
your resources

Semantic
layers on top
of evented I/O

You're using
Threads

You're using
Threads
(in degenerate form)

Exercises

Echo Server

- Listen on a port
- Respond to each request in a new Thread
- **Extra Credit:** Record stats on requests in a shared data structure

Connection Pool

- Allow N clients to access X shared instances of, say, Redis (where $N > X$)
- Clients “check out” a connection and get exclusive access
- Clients “check in” when done
- Instances get re-used



livingsocial

THANKS!

david.copeland@livingsocial.com

[@davetron5000](https://twitter.com/davetron5000)

www.naildrivin5.com

www.awesomecommandlineapps.com